

Ji Sung Han

+1 647-321-9811 | jishan.han@mail.utoronto.ca | [GitHub](#) * | [LinkedIn](#) * | [Personal Website](#) *

EDUCATION

University of Toronto

Sep. 2021 – Dec. 2026

Honours Bachelor of Science in **Computer Science and Statistics**

- Technical Director at UofT Blockchain, Teaching Assistant (Calculus), Focus in AI & Machine Learning.

Relevant Coursework: Databases, Operating Systems, Software Engineering, Algorithm Design, Deep Learning

EXPERIENCE

NVIDIA

Winter 2026

Systems Software Engineer Intern

Toronto, ON

- Improve **system performance and reliability** for **AI/HPC stacks** by profiling and debugging **C++/Linux** services, and building **Python** tooling for benchmarking, telemetry, and regression detection.

Tesla

Sep. 2025 – Jan. 2026

Data Engineer Intern

Palo Alto, CA

- Architected a real-time **Kafka + Spark Streaming** pipeline on **Kubernetes** processing **50M+ events/day** for global operations, reducing data lag from **minutes to seconds** with **auto-scaling** and maintaining **p95 latency <1s**.
- Built and **productionized** a mission-critical batch **ETL + ML** pipeline on **Spark/Airflow** for **inactive installer risk prediction**, driving **~5 account reactivations/week** and delivering an **estimated \$15M/year in value**.
- Optimized **Airflow DAGs** across the risk prediction pipeline processing **1M+ rows/day** via vectorized **pandas**, task parallelization, and **SQL partitioning/indexing**, cutting runtime from **5h to 5m** and enabling daily decision cycles.
- Built and deployed an **XGBoost Top-K model** with rule-based signals and **HITL workflows (FastAPI)** for large-scale entity matching (500K+ entities/day), achieving **97% accuracy** while automating **90% of manual operations**.
- Shipped **CI/CD** workflows with **Docker, Jenkins/GitHub Actions**, and **Kubernetes**, orchestrating **Airflow** jobs with retries, alerting, and **Data Catalog** integration for **global rollouts**, reducing release cycles by **40%**.

Colliers

May 2025 – Aug. 2025

Software Engineer Intern

Toronto, ON

- Developed a **FastAPI/Angular** platform (**Python, TypeScript, HTML/CSS**) on **GCP** with **CI/CD**, enabling **40+ brokers** to analyze leases 3x faster through **semantic search and summarization** powered by **LangChain** and **Vertex AI**.
- Reduced **frontend query latency by 30%** (5.0s to 3.5s) via **Redis caching (GCP Memorystore)** with hashed keys and TTL tuning; enabled **async summarization** using **Python asyncio** and partitioned **BigQuery (SQL)** queries.
- Built a **distributed OCR + LLM-based ELT pipeline** on **Airflow + PySpark (Dataproc)** to extract **60+ fields** from **10k+ PDF/HTML leases**, reducing **hourly batch time by 52%** (58m to 28m) via **batching and repartition tuning**.
- Replaced 1–3 day manual extraction with automated hourly ELT to **GCS (JSON/CSV) → BigQuery (SQL)** feeding Power BI dashboards, enabling **1000+ analysts** to refresh in < **1hr** and saving **\$25k/yr** by eliminating manual entry.

LangChain

Sep. 2024 – Dec. 2024

Software Engineer, Open Source

Remote

- Built a **Dropbox document loader** in **TypeScript/Node.js** for **LangChain**, supporting 10k+ files via recursive traversal and temp-file parsing; integrated with **UnstructuredLoader** for document conversion and metadata enrichment.

Maple Bit Services

May 2022 – Aug. 2022

Software Engineer Intern

Toronto, ON

- Improved POS speed by **35%** via **async FastAPI endpoints** on **AWS EC2**, optimized **MySQL/PostgreSQL (AWS RDS)** joins and indexing, and offloaded RPC load using **C++ gRPC** microservices; cut peak-hour latency by 200ms.

University of Toronto

May 2024 – May 2025

Research Assistant / Technical Director

Toronto, ON

- Built a **transformer ensemble** (BERT/RoBERTa/XLNet) in **PyTorch/TensorFlow** for crypto fraud detection; improved F1 by **12%** with **GNN-augmented encoders** and **Integrated Gradients** token attribution.

PROJECTS | [View All](#) *

Lifomation (Cofounder)

| Angular, Azure OpenAI, Redis, BullMQ, Jenkins, Docker, Dynatrace

- Co-developed a cloud-native AI document platform (8-person team) with **Angular** frontend and **Azure OpenAI**-powered **RAG**; built scalable workflows using **Redis (pub-sub)** and **BullMQ**.
- Enhanced Azure OpenAI privacy features for **1,000+ users**, resulting in a **Sequoia Capital** product demo.

TECHNICAL SKILLS

Languages: Python, JavaScript (TypeScript), C, C++, C#, Java, Kotlin, Swift, R, Bash, Haskell, HTML/CSS, XML

Frameworks: React, NEXT.js, Angular, Spring Boot, Node.js, Express.js, Flask, Django, .NET, TensorFlow, PyTorch

Cloud & DevOps: AWS, GCP, Azure, Docker, Kubernetes, OpenShift, Jenkins, GitHub Actions, gRPC